

In [1]:

```
1 import os
2 import librosa
3 import numpy as np
4 import tensorflow as tf
5 from tensorflow.keras.models import Sequential
6 from tensorflow.keras.layers import LSTM, Dense, Dropout
7 from tensorflow.keras.callbacks import EarlyStopping
8 from sklearn.model_selection import train_test_split
```

C:\Users\shrut\anaconda3\lib\site-packages\numpy_distributor_init.py:30: UserWarning: loaded more than 1 DLL from .libs:
C:\Users\shrut\anaconda3\lib\site-packages\numpy\.libs\libopenblas.GK7GX5KEQ4F6UYO3P26ULGBQYHGQ07J4.gfortran-win_amd64.dll
C:\Users\shrut\anaconda3\lib\site-packages\numpy\.libs\libopenblas.WCDJNK7YVM PZQ2ME2ZZHJJRJ3JIKNDB7.gfortran-win_amd64.dll
warnings.warn("loaded more than 1 DLL from .libs:")
C:\Users\shrut\anaconda3\lib\site-packages\tensorflow\python\keras\engine\training_arrays_v1.py:37: UserWarning: A NumPy version >=1.22.4 and <2.3.0 is required for this version of SciPy (detected version 1.21.0)
from scipy.sparse import issparse # pylint: disable=g-import-not-at-top

In [2]:

```
1 # Set paths
2 DATA_PATH = 'Audio_Files/' # Update this with your path
3 major_folder = os.path.join(DATA_PATH, 'major')
4 minor_folder = os.path.join(DATA_PATH, 'minor')
5
6 # Hyperparameters
7 n_mfcc = 20
8 sequence_length = 30
9 batch_size = 32
10 epochs = 50
```

In [3]:

```
1 # Function to load and preprocess audio files
2 def load_and_extract_features(folder_path):
3     features = []
4     for file_name in os.listdir(folder_path):
5         if file_name.endswith('.wav'):
6             file_path = os.path.join(folder_path, file_name)
7             y, sr = librosa.load(file_path)
8             mfccs = librosa.feature.mfcc(y=y, sr=sr, n_mfcc=n_mfcc)
9             features.append(mfccs.T)
10     return features
```

```
In [4]: 1 # Load and prepare dataset
2 major_features = load_and_extract_features(major_folder)
3 minor_features = load_and_extract_features(minor_folder)
4
5 # Prepare sequences
6 def create_sequences(features, seq_length):
7     X, y = [], []
8     for mfcc in features:
9         for i in range(len(mfcc) - seq_length):
10             X.append(mfcc[i:i + seq_length])
11             y.append(mfcc[i + seq_length])
12     return np.array(X), np.array(y)
```

```
In [5]: 1 X_major, y_major = create_sequences(major_features, sequence_length)
2 X_minor, y_minor = create_sequences(minor_features, sequence_length)
3
4 # Combine and label data (1 for major, 0 for minor)
5 X = np.concatenate((X_major, X_minor), axis=0)
6 y = np.concatenate((y_major, y_minor), axis=0)
```

```
In [6]: 1 # Split data into train and test sets
2 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, r
```

```
In [7]: 1 # Define the model
2 model = Sequential([
3     LSTM(64, input_shape=(sequence_length, n_mfcc), return_sequences=True),
4     Dropout(0.2),
5     LSTM(64),
6     Dropout(0.2),
7     Dense(32, activation='relu'),
8     Dense(n_mfcc) # Output layer with the same dimension as MFCC input
9 ])
```

C:\Users\shrut\anaconda3\lib\site-packages\keras\src\layers\rnn\rnn.py:204: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(**kwargs)
```

```
In [*]: 1 model.compile(optimizer='adam', loss='mse')
2 model.summary()
3
4 # Set up early stopping
5 early_stopping = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True)
6
7 # Train the model with early stopping
8 history = model.fit(
9     X_train, y_train,
10    batch_size=batch_size,
11    epochs=epochs,
12    validation_split=0.2,
13    callbacks=[early_stopping]
14 )
```

Model: "sequential"


Layer (type)	Output Shape	
lstm (LSTM)	(None, 30, 64)	
dropout (Dropout)	(None, 30, 64)	
lstm_1 (LSTM)	(None, 64)	
dropout_1 (Dropout)	(None, 64)	
dense (Dense)	(None, 32)	
dense_1 (Dense)	(None, 20)	





Total params: 57,524 (224.70 KB)


Trainable params: 57,524 (224.70 KB)


Non-trainable params: 0 (0.00 B)


Epoch 1/50
1148/1148  52s 41ms/step - loss: 4055.2861 - val_loss: 60.6.1160


Epoch 2/50
1148/1148  45s 40ms/step - loss: 490.6552 - val_loss: 143.1201


Epoch 3/50
1148/1148  41s 35ms/step - loss: 192.3761 - val_loss: 109.3244


Epoch 4/50
1148/1148  40s 35ms/step - loss: 165.5600 - val_loss: 100.6715


Epoch 5/50
1148/1148  42s 36ms/step - loss: 155.7856 - val_loss: 97.2126


Epoch 6/50
1148/1148  39s 34ms/step - loss: 150.0359 - val_loss: 96.6808


Epoch 7/50
1148/1148  39s 34ms/step - loss: 146.2354 - val_loss: 92.5811


Epoch 8/50
1148/1148  41s 35ms/step - loss: 142.7170 - val_loss: 89.9662


Epoch 9/50
1148/1148  40s 35ms/step - loss: 138.0513 - val_loss: 87.4177


Epoch 10/50
1148/1148  42s 37ms/step - loss: 134.7647 - val_loss: 87.6687


Epoch 11/50
1148/1148  83s 37ms/step - loss: 129.9500 - val_loss: 83.6419


Epoch 12/50
1148/1148  41s 36ms/step - loss: 130.1402 - val_loss: 79.7878


Epoch 13/50
1148/1148  42s 36ms/step - loss: 122.5917 - val_loss: 81.1944


Epoch 14/50
1148/1148  40s 35ms/step - loss: 120.8897 - val_loss: 77.4183














Epoch 15/50
1148/1148  41s 36ms/step - loss: 118.4245 - val_loss: 75.8484

Epoch 16/50
1148/1148  45s 39ms/step - loss: 114.2781 - val_loss: 75.6577

Epoch 17/50
1148/1148  42s 37ms/step - loss: 112.3637 - val_loss: 73.7750

Epoch 18/50
1148/1148  44s 38ms/step - loss: 108.2301 - val_loss: 71.2567

Epoch 19/50
1148/1148  46s 40ms/step - loss: 104.5000 - val_loss: 72.1390

Epoch 20/50
1148/1148  **52s** 45ms/step - loss: 100.4783 - val_loss: 67.2329
Epoch 21/50
1148/1148  **129s** 86ms/step - loss: 95.1257 - val_loss: 62.8173
Epoch 22/50
1148/1148  **148s** 91ms/step - loss: 90.4192 - val_loss: 59.2195
Epoch 23/50
1148/1148  **139s** 88ms/step - loss: 85.2510 - val_loss: 59.1313
Epoch 24/50
1148/1148  **129s** 76ms/step - loss: 80.3584 - val_loss: 53.4864
Epoch 25/50
1148/1148  **57s** 49ms/step - loss: 75.3179 - val_loss: 51.2329
Epoch 26/50
1148/1148  **101s** 88ms/step - loss: 70.8768 - val_loss: 49.9078
Epoch 27/50
1148/1148  **73s** 63ms/step - loss: 67.9127 - val_loss: 47.8595
Epoch 28/50
1148/1148  **80s** 61ms/step - loss: 65.4084 - val_loss: 46.7697
Epoch 29/50
1148/1148  **89s** 78ms/step - loss: 61.8384 - val_loss: 45.6907
Epoch 30/50
1148/1148  **150s** 84ms/step - loss: 60.6193 - val_loss: 43.2519
Epoch 31/50
1148/1148  **135s** 78ms/step - loss: 57.6541 - val_loss: 41.1966
Epoch 32/50
1134/1148  **1s** 88ms/step - loss: 56.9344

```
In [*]: 1 # Evaluate
2 test_loss = model.evaluate(X_test, y_test)
3 print(f"Test Loss (MSE): {test_loss}")
4
5 # Save model
6 model.save("melody_generator_model.h5")
```

```
In [*]: 1 # Training Loss (MSE) on the training data
2 train_loss = model.evaluate(X_train, y_train)
3 print(f"Training Loss (MSE): {train_loss}")
4
5 # Calculate test Loss (MSE) on the test data
6 test_loss = model.evaluate(X_test, y_test)
7 print(f"Test Loss (MSE): {test_loss}")
```

```
In [*]: 1 from sklearn.metrics import r2_score
2
3 # Predict on train and test sets
4 y_train_pred = model.predict(X_train)
5 y_test_pred = model.predict(X_test)
6
7 # Calculate R-squared score for training and test sets
8 train_r2 = r2_score(y_train, y_train_pred)
9 test_r2 = r2_score(y_test, y_test_pred)
10
11 print(f"Training R^2 Score: {train_r2}")
12 print(f"Test R^2 Score: {test_r2}")
```

```
In [*]: 1 import librosa.display
2 import matplotlib.pyplot as plt
3
4 # Function to generate a sequence
5 def generate_sequence(model, start_sequence, sequence_length):
6     generated_sequence = start_sequence.copy()
7
8     for _ in range(sequence_length):
9         prediction = model.predict(generated_sequence[-30:].reshape(1, -1))
10        generated_sequence = np.vstack([generated_sequence, prediction])
11
12    return generated_sequence
```

```
In [*]: 1 def mfcc_to_audio(mfcc_sequence, sr=22050):
2     S = librosa.feature.inverse.mfcc_to_mel(mfcc_sequence.T, n_mels=128)
3     audio = librosa.feature.inverse.mel_to_audio(S, sr=sr)
4     return audio
```

```
In [*]: 1 import IPython.display as ipd
2
3 # Load a seed audio file and extract initial MFCCs
4 y, sr = librosa.load(os.path.join(major_folder, "Major_8.wav")) # replace
5 initial_mfccs = librosa.feature.mfcc(y=y, sr=sr, n_mfcc=n_mfcc).T[:sequence_length]
6
7 # Generating melody
8 generated_mfccs = generate_sequence(model, initial_mfccs, sequence_length=sequence_length)
9
10 # Convert generated MFCCs to audio
11 generated_audio = mfcc_to_audio(generated_mfccs, sr=sr)
12
13 # Play generated audio
14 ipd.Audio(generated_audio, rate=sr)
15
```

```
In [*]: 1 import IPython.display as ipd
2
3 # Load a seed audio file and extract initial MFCCs
4 y, sr = librosa.load(os.path.join(minor_folder, "Minor_8.wav")) # replace
5 initial_mfccs = librosa.feature.mfcc(y=y, sr=sr, n_mfcc=n_mfcc).T[:sequence_length]
6
7 # Generate a melody
8 generated_mfccs = generate_sequence(model, initial_mfccs, sequence_length=sequence_length)
9
10 # Convert generated MFCCs to audio
11 generated_audio = mfcc_to_audio(generated_mfccs, sr=sr)
12
13 # Play generated audio
14 ipd.Audio(generated_audio, rate=sr)
15
```

```
In [*]: 1 import librosa.display
2 import numpy as np
3 import matplotlib.pyplot as plt
4
5 # Function to generate a sequence with added Gaussian noise
6 def generate_sequence_with_noise(model, start_sequence, sequence_length, sr):
7     generated_sequence = start_sequence.copy()
8
9     for _ in range(sequence_length):
10         # Predict the next MFCC frame
11         prediction = model.predict(generated_sequence[-30:].reshape(1, -1), sr=sr)
12
13         # Add Gaussian noise to the prediction for variability
14         prediction += np.random.normal(0, noise_stddev, prediction.shape)
15
16         # Append the predicted frame to the generated sequence
17         generated_sequence = np.vstack([generated_sequence, prediction])
18
19     return generated_sequence
```

```
In [*]: 1 # Function to convert MFCC back to audio
2 def mfcc_to_audio(mfcc_sequence, sr=22050):
3     S = librosa.feature.inverse.mfcc_to_mel(mfcc_sequence.T, n_mels=128)
4     audio = librosa.feature.inverse.mel_to_audio(S, sr=sr)
5     return audio
6
7 # Load a seed audio file and extract MFCCs
8 y, sr = librosa.load(os.path.join(major_folder, "Major_9.wav")) # replace
9 initial_mfccs = librosa.feature.mfcc(y=y, sr=sr, n_mfcc=n_mfcc).T[:sequence_length]
10
```

```
In [*]: 1 # Generate the melody with Gaussian noise added
2 generated_mfccs = generate_sequence_with_noise(model, initial_mfccs, sequence_length, sr)
3
4 # Convert generated MFCCs to audio
5 generated_audio = mfcc_to_audio(generated_mfccs, sr=sr)
```

```
In [*]: 1 # Plot the spectrogram of the generated audio
2 plt.figure(figsize=(10, 6))
3 librosa.display.waveshow(generated_audio, sr=sr)
4 plt.title("Generated Audio Waveform")
5 plt.show()
6
7 plt.figure(figsize = (10,6))
8 librosa.display.waveshow(y = y, sr = sr)
9 plt.title("Original Audio Waveform")
10 plt.show()
```

```
In [*]: 1 # Play the generated audio
2 import IPython.display as ipd
3 ipd.Audio(generated_audio, rate=sr)
```